

Activity_Course 5 Waze project lab

March 3, 2025

1 Waze Project

Course 5 - Regression analysis: Simplify complex data relationships

Your team is more than halfway through their user churn project. Earlier, you completed a project proposal, used Python to explore and analyze Waze's user data, created data visualizations, and conducted a hypothesis test. Now, leadership wants your team to build a regression model to predict user churn based on a variety of variables.

You check your inbox and discover a new email from Ursula Sayo, Waze's Operations Manager. Ursula asks your team about the details of the regression model. You also notice two follow-up emails from your supervisor, May Santner. The first email is a response to Ursula, and says that the team will build a binomial logistic regression model. In her second email, May asks you to help build the model and prepare an executive summary to share your results.

A notebook was structured and prepared to help you in this project. Please complete the following questions and prepare an executive summary.

2 Course 5 End-of-course project: Regression modeling

In this activity, you will build a binomial logistic regression model. As you have learned, logistic regression helps you estimate the probability of an outcome. For data science professionals, this is a useful skill because it allows you to consider more than one variable against the variable you're measuring against. This opens the door for much more thorough and flexible analysis to be completed.

The purpose of this project is to demonstrate knowledge of exploratory data analysis (EDA) and a binomial logistic regression model.

The goal is to build a binomial logistic regression model and evaluate the model's performance.

This activity has three parts:

Part 1: EDA & Checking Model Assumptions * What are some purposes of EDA before constructing a binomial logistic regression model?

Part 2: Model Building and Evaluation * What resources do you find yourself using as you complete this stage?

Part 3: Interpreting Model Results

- What key insights emerged from your model(s)?
- What business recommendations do you propose based on the models built?

Follow the instructions and answer the question below to complete the activity. Then, you will complete an executive summary using the questions listed on the PACE Strategy Document.

Be sure to complete this activity before moving on. The next course item will provide you with a completed exemplar to compare to your own work.

3 Build a regression model

4 PACE stages

Throughout these project notebooks, you'll see references to the problem-solving framework PACE. The following notebook components are labeled with the respective PACE stage: Plan, Analyze, Construct, and Execute.

4.1 PACE: Plan

Consider the questions in your PACE Strategy Document to reflect on the Plan stage.

4.1.1 Task 1. Imports and data loading

Import the data and packages that you've learned are needed for building logistic regression models.

```
[1]: # Packages for numerics + dataframes
import pandas as pd
import numpy as np

# Packages for visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Packages for Logistic Regression & Confusion Matrix
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, accuracy_score, \
    precision_score, \
    recall_score, f1_score, confusion_matrix, ConfusionMatrixDisplay
from sklearn.linear_model import LogisticRegression
```

Import the dataset.

Note: As shown in this cell, the dataset has been automatically loaded in for you. You do not need to download the .csv file, or provide more code, in order to access the dataset and proceed with this lab. Please continue with this activity by completing the following instructions.

```
[2]: # Load the dataset by running this cell

df = pd.read_csv('waze_dataset.csv')
```

4.2 PACE: Analyze

Consider the questions in your PACE Strategy Document to reflect on the Analyze stage.

In this stage, consider the following question:

- What are some purposes of EDA before constructing a binomial logistic regression model?

==> ENTER YOUR RESPONSE HERE

4.2.1 Task 2a. Explore data with EDA

Analyze and discover data, looking for correlations, missing data, potential outliers, and/or duplicates.

Start with `.shape` and `info()`.

```
[3]: ### YOUR CODE HERE ###
print(df.shape)
df.info
```

```
(14999, 13)
```

```
[3]: <bound method DataFrame.info of          ID      label  sessions  drives
total_sessions \
```

```
0          0  retained      283    226    296.748273
1          1  retained      133    107    326.896596
2          2  retained      114     95    135.522926
3          3  retained       49     40     67.589221
4          4  retained       84     68    168.247020
```

```
...      ...      ...      ...      ...
14994  14994  retained       60     55    207.875622
14995  14995  retained       42     35    187.670313
14996  14996  retained      273    219    422.017241
14997  14997   churned      149    120    180.524184
14998  14998  retained       73     58    353.419797
```

```
          n_days_after_onboarding  total_navigations_fav1 \
0                               2276                    208
1                               1225                     19
2                               2651                     0
3                               15                      322
4                               1562                    166
```

```
...      ...      ...
```

14994	140	317
14995	2505	15
14996	1873	17
14997	3150	45
14998	3383	13

	total_navigations_fav2	driven_km_drives	duration_minutes_drives \
0	0	2628.845068	1985.775061
1	64	13715.920550	3160.472914
2	0	3059.148818	1610.735904
3	7	913.591123	587.196542
4	5	3950.202008	1219.555924
...
14994	0	2890.496901	2186.155708
14995	10	4062.575194	1208.583193
14996	0	3097.825028	1031.278706
14997	0	4051.758549	254.187763
14998	51	6030.498773	3042.436423

	activity_days	driving_days	device
0	28	19	Android
1	13	11	iPhone
2	14	8	Android
3	7	3	iPhone
4	27	18	Android
...
14994	25	17	iPhone
14995	25	20	Android
14996	18	17	iPhone
14997	6	6	iPhone
14998	14	13	iPhone

[14999 rows x 13 columns]>

Question: Are there any missing values in your data?

==> ENTER YOUR RESPONSE HERE

Use `.head()`.

```
[4]: ### YOUR CODE HERE ###
df.head
```

```
[4]: <bound method NDFrame.head of          ID    label  sessions  drives
total_sessions \
0          0  retained    283    226    296.748273
1          1  retained    133    107    326.896596
2          2  retained    114     95    135.522926
```

3	3	retained	49	40	67.589221
4	4	retained	84	68	168.247020
...
14994	14994	retained	60	55	207.875622
14995	14995	retained	42	35	187.670313
14996	14996	retained	273	219	422.017241
14997	14997	churned	149	120	180.524184
14998	14998	retained	73	58	353.419797

	n_days_after_onboarding	total_navigations_fav1 \
0	2276	208
1	1225	19
2	2651	0
3	15	322
4	1562	166
...
14994	140	317
14995	2505	15
14996	1873	17
14997	3150	45
14998	3383	13

	total_navigations_fav2	driven_km_drives	duration_minutes_drives \
0	0	2628.845068	1985.775061
1	64	13715.920550	3160.472914
2	0	3059.148818	1610.735904
3	7	913.591123	587.196542
4	5	3950.202008	1219.555924
...
14994	0	2890.496901	2186.155708
14995	10	4062.575194	1208.583193
14996	0	3097.825028	1031.278706
14997	0	4051.758549	254.187763
14998	51	6030.498773	3042.436423

	activity_days	driving_days	device
0	28	19	Android
1	13	11	iPhone
2	14	8	Android
3	7	3	iPhone
4	27	18	Android
...
14994	25	17	iPhone
14995	25	20	Android
14996	18	17	iPhone
14997	6	6	iPhone
14998	14	13	iPhone

[14999 rows x 13 columns]>

Use `.drop()` to remove the ID column since we don't need this information for your analysis.

```
[5]: ### YOUR CODE HERE ###
df = df.drop('ID', axis=1)
```

Now, check the class balance of the dependent (target) variable, `label`.

```
[6]: ### YOUR CODE HERE ###
df['label'].value_counts(normalize=True)
```

```
[6]: retained    0.822645
churned        0.177355
Name: label, dtype: float64
```

Call `.describe()` on the data.

```
[7]: ### YOUR CODE HERE ###
df.describe()
```

```
[7]:
```

	sessions	drives	total_sessions	n_days_after_onboarding \
count	14999.000000	14999.000000	14999.000000	14999.000000
mean	80.633776	67.281152	189.964447	1749.837789
std	80.699065	65.913872	136.405128	1008.513876
min	0.000000	0.000000	0.220211	4.000000
25%	23.000000	20.000000	90.661156	878.000000
50%	56.000000	48.000000	159.568115	1741.000000
75%	112.000000	93.000000	254.192341	2623.500000
max	743.000000	596.000000	1216.154633	3500.000000

	total_navigations_fav1	total_navigations_fav2	driven_km_drives \
count	14999.000000	14999.000000	14999.000000
mean	121.605974	29.672512	4039.340921
std	148.121544	45.394651	2502.149334
min	0.000000	0.000000	60.441250
25%	9.000000	0.000000	2212.600607
50%	71.000000	9.000000	3493.858085
75%	178.000000	43.000000	5289.861262
max	1236.000000	415.000000	21183.401890

	duration_minutes_drives	activity_days	driving_days
count	14999.000000	14999.000000	14999.000000
mean	1860.976012	15.537102	12.179879
std	1446.702288	9.004655	7.824036
min	18.282082	0.000000	0.000000
25%	835.996260	8.000000	5.000000

50%	1478.249859	16.000000	12.000000
75%	2464.362632	23.000000	19.000000
max	15851.727160	31.000000	30.000000

Question: Are there any variables that could potentially have outliers just by assessing at the quartile values, standard deviation, and max values?

==> ENTER YOUR RESPONSE HERE

4.2.2 Task 2b. Create features

Create features that may be of interest to the stakeholder and/or that are needed to address the business scenario/problem.

km_per_driving_day You know from earlier EDA that churn rate correlates with distance driven per driving day in the last month. It might be helpful to engineer a feature that captures this information.

1. Create a new column in `df` called `km_per_driving_day`, which represents the mean distance driven per driving day for each user.
2. Call the `describe()` method on the new column.

```
[8]: # 1. Create `km_per_driving_day` column
df['km_per_driving_day'] = df['driven_km_drives'] / df['driving_days']

# 2. Call `describe()` on the new column
df['km_per_driving_day'].describe()
```

```
[8]: count    1.499900e+04
mean              inf
std              NaN
min    3.022063e+00
25%    1.672804e+02
50%    3.231459e+02
75%    7.579257e+02
max              inf
Name: km_per_driving_day, dtype: float64
```

Note that some values are infinite. This is the result of there being values of zero in the `driving_days` column. Pandas imputes a value of infinity in the corresponding rows of the new column because division by zero is undefined.

1. Convert these values from infinity to zero. You can use `np.inf` to refer to a value of infinity.
2. Call `describe()` on the `km_per_driving_day` column to verify that it worked.

```
[9]: # 1. Convert infinite values to zero
df.loc[df['km_per_driving_day']==np.inf, 'km_per_driving_day'] = 0
```

```
# 2. Confirm that it worked
df['km_per_driving_day'].describe()
```

```
[9]: count    14999.000000
     mean      578.963113
     std      1030.094384
     min        0.000000
     25%      136.238895
     50%      272.889272
     75%      558.686918
     max     15420.234110
     Name: km_per_driving_day, dtype: float64
```

professional_driver Create a new, binary feature called **professional_driver** that is a 1 for users who had 60 or more drives **and** drove on 15+ days in the last month.

Note: The objective is to create a new feature that separates professional drivers from other drivers. In this scenario, domain knowledge and intuition are used to determine these deciding thresholds, but ultimately they are arbitrary.

To create this column, use the `np.where()` function. This function accepts as arguments: 1. A condition 2. What to return when the condition is true 3. What to return when the condition is false

Example:

```
x = [1, 2, 3]
x = np.where(x > 2, 100, 0)
x
array([ 0,  0, 100])
```

```
[10]: # Create `professional_driver` column
df['professional_driver'] = np.where((df['drives'] >= 60) & (df['driving_days'] >= 15), 1, 0)
```

Perform a quick inspection of the new variable.

1. Check the count of professional drivers and non-professionals
2. Within each class (professional and non-professional) calculate the churn rate

```
[11]: # 1. Check count of professionals and non-professionals
print(df['professional_driver'].value_counts())

# 2. Check in-class churn rate
df.groupby(['professional_driver'])['label'].value_counts(normalize=True)
```

```
0    12405
1     2594
```



```
Name: professional_driver, dtype: int64
```

```
[11]: professional_driver  label
      0                  retained    0.801202
      1                  churned    0.198798
      1                  retained    0.924437
      1                  churned    0.075563
      Name: label, dtype: float64
```

The churn rate for professional drivers is 7.6%, while the churn rate for non-professionals is 19.9%. This seems like it could add predictive signal to the model.

4.3 PACE: Construct

After analysis and deriving variables with close relationships, it is time to begin constructing the model.

Consider the questions in your PACE Strategy Document to reflect on the Construct stage.

In this stage, consider the following question:

- Why did you select the X variables you did?

==> ENTER YOUR RESPONSE HERE

4.3.1 Task 3a. Preparing variables

Call `info()` on the dataframe to check the data type of the `label` variable and to verify if there are any missing values.

```
[12]: ### YOUR CODE HERE ###
      df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
Data columns (total 14 columns):
 #   Column                                Non-Null Count  Dtype  
---  -
 0   label                                14299 non-null  object  
 1   sessions                             14999 non-null  int64   
 2   drives                               14999 non-null  int64   
 3   total_sessions                       14999 non-null  float64  
 4   n_days_after_onboarding              14999 non-null  int64   
 5   total_navigations_fav1               14999 non-null  int64   
 6   total_navigations_fav2               14999 non-null  int64   
 7   driven_km_drives                     14999 non-null  float64  
 8   duration_minutes_drives               14999 non-null  float64  
 9   activity_days                         14999 non-null  int64   
10  driving_days                          14999 non-null  int64   
```

```

11 device                14999 non-null object
12 km_per_driving_day    14999 non-null float64
13 professional_driver    14999 non-null int64
dtypes: float64(4), int64(8), object(2)
memory usage: 1.6+ MB

```

Because you know from previous EDA that there is no evidence of a non-random cause of the 700 missing values in the `label` column, and because these observations comprise less than 5% of the data, use the `dropna()` method to drop the rows that are missing this data.

```

[14]: # Drop rows with missing data in `label` column
      ### YOUR CODE HERE ###
      df = df.dropna(subset=['label'])

```

Impute outliers You rarely want to drop outliers, and generally will not do so unless there is a clear reason for it (e.g., typographic errors).

At times outliers can be changed to the **median, mean, 95th percentile, etc.**

Previously, you determined that seven of the variables had clear signs of containing outliers:

- `sessions`
- `drives`
- `total_sessions`
- `total_navigations_fav1`
- `total_navigations_fav2`
- `driven_km_drives`
- `duration_minutes_drives`

For this analysis, impute the outlying values for these columns. Calculate the **95th percentile** of each column and change to this value any value in the column that exceeds it.

```

[15]: # Impute outliers
      for column in ['sessions', 'drives', 'total_sessions', 'total_navigations_fav1',
                    'total_navigations_fav2', 'driven_km_drives',
                    'duration_minutes_drives']:
          threshold = df[column].quantile(0.95)
          df.loc[df[column] > threshold, column] = threshold

```

Call `describe()`.

```

[16]: ### YOUR CODE HERE ###
      df.describe()

```

```

[16]:
      sessions      drives  total_sessions  n_days_after_onboarding \
count  14299.000000  14299.000000    14299.000000             14299.000000
mean       76.539688    63.964683     183.717304             1751.822505
std       67.243178    55.127927     118.720520             1008.663834
min        0.000000     0.000000      0.220211              4.000000
25%       23.000000    20.000000     90.457733             878.500000

```

50%	56.000000	48.000000	158.718571	1749.000000
75%	111.000000	93.000000	253.540450	2627.500000
max	243.000000	200.000000	455.439492	3500.000000

	total_navigations_fav1	total_navigations_fav2	driven_km_drives \
count	14299.000000	14299.000000	14299.000000
mean	114.562767	27.187216	3944.558631
std	124.378550	36.715302	2218.358258
min	0.000000	0.000000	60.441250
25%	10.000000	0.000000	2217.319909
50%	71.000000	9.000000	3496.545617
75%	178.000000	43.000000	5299.972162
max	422.000000	124.000000	8898.716275

	duration_minutes_drives	activity_days	driving_days \
count	14299.000000	14299.000000	14299.000000
mean	1792.911210	15.544653	12.182530
std	1224.329759	9.016088	7.833835
min	18.282082	0.000000	0.000000
25%	840.181344	8.000000	5.000000
50%	1479.394387	16.000000	12.000000
75%	2466.928876	23.000000	19.000000
max	4668.180092	31.000000	30.000000

	km_per_driving_day	professional_driver
count	14299.000000	14299.000000
mean	581.942399	0.173998
std	1038.254509	0.379121
min	0.000000	0.000000
25%	136.168003	0.000000
50%	273.301012	0.000000
75%	558.018761	0.000000
max	15420.234110	1.000000

Encode categorical variables Change the data type of the `label` column to be binary. This change is needed to train a logistic regression model.

Assign a 0 for all `retained` users.

Assign a 1 for all `churned` users.

Save this variable as `label2` as to not overwrite the original `label` variable.

Note: There are many ways to do this. Consider using `np.where()` as you did earlier in this notebook.

```
[17]: # Create binary `label2` column
df['label2'] = np.where(df['label']=='churned', 1, 0)
```

```
df[['label', 'label2']].tail()
```

```
[17]:      label  label2
14994  retained      0
14995  retained      0
14996  retained      0
14997   churned      1
14998  retained      0
```

4.3.2 Task 3b. Determine whether assumptions have been met

The following are the assumptions for logistic regression:

- Independent observations (This refers to how the data was collected.)
- No extreme outliers
- Little to no multicollinearity among X predictors
- Linear relationship between X and the **logit** of y

For the first assumption, you can assume that observations are independent for this project.

The second assumption has already been addressed.

The last assumption will be verified after modeling.

Note: In practice, modeling assumptions are often violated, and depending on the specifics of your use case and the severity of the violation, it might not affect your model much at all or it will result in a failed model.

Collinearity Check the correlation among predictor variables. First, generate a correlation matrix.

```
[19]: # Generate a correlation matrix
      ### YOUR CODE HERE ###
      df.corr(method='pearson')
```

```
[19]:      sessions    drives  total_sessions  \
sessions      1.000000  0.996942      0.597189
drives         0.996942  1.000000      0.595285
total_sessions  0.597189  0.595285      1.000000
n_days_after_onboarding  0.007101  0.006940      0.006596
total_navigations_fav1  0.001858  0.001058      0.000187
total_navigations_fav2  0.008536  0.009505      0.010371
driven_km_drives      0.002996  0.003445      0.001016
duration_minutes_drives -0.004545 -0.003889     -0.000338
activity_days      0.025113  0.024357      0.015755
driving_days       0.020294  0.019608      0.012953
km_per_driving_day  -0.011569 -0.010989     -0.016167
```

professional_driver	0.443654	0.444425	0.254433
label2	0.034911	0.035865	0.024568

	n_days_after_onboarding	total_navigations_fav1	\
sessions	0.007101	0.001858	
drives	0.006940	0.001058	
total_sessions	0.006596	0.000187	
n_days_after_onboarding	1.000000	-0.002450	
total_navigations_fav1	-0.002450	1.000000	
total_navigations_fav2	-0.004968	0.002866	
driven_km_drives	-0.004652	-0.007368	
duration_minutes_drives	-0.010167	0.005646	
activity_days	-0.009418	0.010902	
driving_days	-0.007321	0.010419	
km_per_driving_day	0.011764	-0.000197	
professional_driver	0.003770	-0.000224	
label2	-0.129263	0.052322	

	total_navigations_fav2	driven_km_drives	\
sessions	0.008536	0.002996	
drives	0.009505	0.003445	
total_sessions	0.010371	0.001016	
n_days_after_onboarding	-0.004968	-0.004652	
total_navigations_fav1	0.002866	-0.007368	
total_navigations_fav2	1.000000	0.003559	
driven_km_drives	0.003559	1.000000	
duration_minutes_drives	-0.003009	0.690515	
activity_days	-0.004425	-0.007441	
driving_days	0.002000	-0.009549	
km_per_driving_day	0.006751	0.344811	
professional_driver	0.007126	-0.000904	
label2	0.015032	0.019767	

	duration_minutes_drives	activity_days	driving_days	\
sessions	-0.004545	0.025113	0.020294	
drives	-0.003889	0.024357	0.019608	
total_sessions	-0.000338	0.015755	0.012953	
n_days_after_onboarding	-0.010167	-0.009418	-0.007321	
total_navigations_fav1	0.005646	0.010902	0.010419	
total_navigations_fav2	-0.003009	-0.004425	0.002000	
driven_km_drives	0.690515	-0.007441	-0.009549	
duration_minutes_drives	1.000000	-0.007895	-0.009425	
activity_days	-0.007895	1.000000	0.947687	
driving_days	-0.009425	0.947687	1.000000	
km_per_driving_day	0.239627	-0.397433	-0.407917	
professional_driver	-0.012128	0.453825	0.469776	
label2	0.040407	-0.303851	-0.294259	

	km_per_driving_day	professional_driver	label2
sessions	-0.011569	0.443654	0.034911
drives	-0.010989	0.444425	0.035865
total_sessions	-0.016167	0.254433	0.024568
n_days_after_onboarding	0.011764	0.003770	-0.129263
total_navigations_fav1	-0.000197	-0.000224	0.052322
total_navigations_fav2	0.006751	0.007126	0.015032
driven_km_drives	0.344811	-0.000904	0.019767
duration_minutes_drives	0.239627	-0.012128	0.040407
activity_days	-0.397433	0.453825	-0.303851
driving_days	-0.407917	0.469776	-0.294259
km_per_driving_day	1.000000	-0.165966	0.148583
professional_driver	-0.165966	1.000000	-0.122312
label2	0.148583	-0.122312	1.000000

Now, plot a correlation heatmap.

```
[20]: # Plot correlation heatmap
plt.figure(figsize=(15,10))
sns.heatmap(df.corr(method='pearson'), vmin=-1, vmax=1, annot=True,
            cmap='coolwarm')
plt.title('Correlation heatmap indicates many low correlated variables',
          fontsize=18)
plt.show();
```



If there are predictor variables that have a Pearson correlation coefficient value greater than the **absolute value of 0.7**, these variables are strongly multicollinear. Therefore, only one of these variables should be used in your model.

Note: 0.7 is an arbitrary threshold. Some industries may use 0.6, 0.8, etc.

Question: Which variables are multicollinear with each other?

==> ENTER YOUR RESPONSE HERE

4.3.3 Task 3c. Create dummies (if necessary)

If you have selected **device** as an X variable, you will need to create dummy variables since this variable is categorical.

In cases with many categorical variables, you can use pandas built-in `pd.get_dummies()`, or you can use scikit-learn's `OneHotEncoder()` function.

Note: Variables with many categories should only be dummied if absolutely necessary. Each category will result in a coefficient for your model which can lead to overfitting.

Because this dataset only has one remaining categorical feature (**device**), it's not necessary to use one of these special functions. You can just implement the transformation directly.

Create a new, binary column called `device2` that encodes user devices as follows:

- Android -> 0
- iPhone -> 1

```
[21]: # Create new `device2` variable
df['device2'] = np.where(df['device']=='Android', 0, 1)
df[['device', 'device2']].tail()
```

```
[21]:      device  device2
14994  iPhone         1
14995  Android         0
14996  iPhone         1
14997  iPhone         1
14998  iPhone         1
```

4.3.4 Task 3d. Model building

Assign predictor variables and target To build your model you need to determine what X variables you want to include in your model to predict your target—`label2`.

Drop the following variables and assign the results to `X`:

- `label1` (this is the target)
- `label2` (this is the target)
- `device` (this is the non-binary-encoded categorical variable)
- `sessions` (this had high multicollinearity)
- `driving_days` (this had high multicollinearity)

Note: Notice that `sessions` and `driving_days` were selected to be dropped, rather than `drives` and `activity_days`. The reason for this is that the features that were kept for modeling had slightly stronger correlations with the target variable than the features that were dropped.

```
[22]: # Isolate predictor variables
X = df.drop(columns = ['label1', 'label2', 'device', 'sessions', 'driving_days'])
```

Now, isolate the dependent (target) variable. Assign it to a variable called `y`.

```
[23]: # Isolate target variable
### YOUR CODE HERE ###
y= df['label2']
```

Split the data Use scikit-learn's `train_test_split()` function to perform a train/test split on your data using the `X` and `y` variables you assigned above.

Note 1: It is important to do a train test to obtain accurate predictions. You always want to fit your model on your training set and evaluate your model on your test set to avoid data leakage.

Note 2: Because the target class is imbalanced (82% retained vs. 18% churned), you want to make sure that you don't get an unlucky split that over- or under-represents the frequency of the minority class. Set the function's `stratify` parameter to `y` to ensure that the minority class appears in both train and test sets in the same proportion that it does in the overall dataset.

```
[24]: # Perform the train-test split
      ### YOUR CODE HERE ###
      # Perform the train-test split
      X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y,
      ↪random_state=42)
```

```
[25]: # Use .head()
      ### YOUR CODE HERE ###
      X_train.head()
```

```
[25]:      drives  total_sessions  n_days_after_onboarding  \
152      108      186.192746      3116
11899     2       3.487590      794
10937    139     347.106403      331
669      108     455.439492     2320
8406     10      89.475821     2478

      total_navigations_fav1  total_navigations_fav2  driven_km_drives  \
152              243              124      8898.716275
11899             114              18      3286.545691
10937              4              7      7400.838975
669              11              4      6566.424830
8406             135              0      1271.248661

      duration_minutes_drives  activity_days  km_per_driving_day  \
152          4668.180092          24      612.305861
11899         1780.902733           5      3286.545691
10937         2349.305267          15      616.736581
669         4558.459870          18      410.401552
8406          938.711572          27       74.779333

      professional_driver  device2
152              1          1
11899             0          1
10937             0          0
669              1          1
8406             0          1
```

Use scikit-learn to instantiate a logistic regression model. Add the argument `penalty = None`.

It is important to add `penalty = None` since your predictors are unscaled.

Refer to scikit-learn's [logistic regression](#) documentation for more information.

Fit the model on `X_train` and `y_train`.

```
[26]: ### YOUR CODE HERE ###
model = LogisticRegression(penalty='none', max_iter=400)

model.fit(X_train, y_train)
```

```
[26]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
        intercept_scaling=1, l1_ratio=None, max_iter=400,
        multi_class='auto', n_jobs=None, penalty='none',
        random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
        warm_start=False)
```

Call the `.coef_` attribute on the model to get the coefficients of each variable. The coefficients are in order of how the variables are listed in the dataset. Remember that the coefficients represent the change in the **log odds** of the target variable for **every one unit increase in X**.

If you want, create a series whose index is the column names and whose values are the coefficients in `model.coef_`.

```
[27]: ### YOUR CODE HERE ###
pd.Series(model.coef_[0], index=X.columns)
```

```
[27]: drives                0.001913
total_sessions           0.000327
n_days_after_onboarding -0.000406
total_navigations_fav1   0.001232
total_navigations_fav2   0.000931
driven_km_drives         -0.000015
duration_minutes_drives   0.000109
activity_days            -0.106032
km_per_driving_day        0.000018
professional_driver       -0.001529
device2                  -0.001041
dtype: float64
```

Call the model's `intercept_` attribute to get the intercept of the model.

```
[28]: ### YOUR CODE HERE ###
model.intercept_
```

```
[28]: array([-0.00170675])
```

Check final assumption Verify the linear relationship between `X` and the estimated log odds (known as logits) by making a regplot.

Call the model's `predict_proba()` method to generate the probability of response for each sample in the training data. (The training data is the argument to the method.) Assign the result to a variable called `training_probabilities`. This results in a 2-D array where each row represents

a user in `X_train`. The first column is the probability of the user not churning, and the second column is the probability of the user churning.

```
[29]: # Get the predicted probabilities of the training data
      ### YOUR CODE HERE ###
      # Get the predicted probabilities of the training data
      training_probabilities = model.predict_proba(X_train)
      training_probabilities
```

```
[29]: array([[0.93963483, 0.06036517],
             [0.61967304, 0.38032696],
             [0.76463181, 0.23536819],
             ...,
             [0.91909641, 0.08090359],
             [0.85092112, 0.14907888],
             [0.93516293, 0.06483707]])
```

In logistic regression, the relationship between a predictor variable and the dependent variable does not need to be linear, however, the log-odds (a.k.a., logit) of the dependent variable with respect to the predictor variable should be linear. Here is the formula for calculating log-odds, where p is the probability of response:

$$\text{logit}(p) = \ln\left(\frac{p}{1-p}\right)$$

1. Create a dataframe called `logit_data` that is a copy of `df`.
2. Create a new column called `logit` in the `logit_data` dataframe. The data in this column should represent the logit for each user.


```
[30]: # 1. Copy the `X_train` dataframe and assign to `logit_data`
      logit_data = X_train.copy()

      # 2. Create a new `logit` column in the `logit_data` df
      logit_data['logit'] = [np.log(prob[1] / prob[0]) for prob in
      ↪ training_probabilities]
```

Plot a regplot where the x-axis represents an independent variable and the y-axis represents the log-odds of the predicted probabilities.

In an exhaustive analysis, this would be plotted for each continuous or discrete predictor variable. Here we show only `driving_days`.

```
[31]: # Plot regplot of `activity_days` log-odds
      ### YOUR CODE HERE ###
      # Plot regplot of `activity_days` log-odds
      sns.regplot(x='activity_days', y='logit', data=logit_data, scatter_kws={'s': 2,
      ↪ 'alpha': 0.5})
      plt.title('Log-odds: activity_days');
```



4.4 PACE: Execute

Consider the questions in your PACE Strategy Document to reflect on the Execute stage.

4.4.1 Task 4a. Results and evaluation

If the logistic assumptions are met, the model results can be appropriately interpreted.

Use the code block below to make predictions on the test data.

```
[35]: # Generate predictions on X_test
      ### YOUR CODE HERE ###
      # Generate predictions on X_test
      y_preds = model.predict(X_test)
```

Now, use the `score()` method on the model with `X_test` and `y_test` as its two arguments. The default score in scikit-learn is **accuracy**. What is the accuracy of your model?

Consider: Is accuracy the best metric to use to evaluate this model?

```
[34]: # Score the model (accuracy) on the test data
      ### YOUR CODE HERE ###
      model.score(X_test, y_test)
```

```
[34]: 0.8237762237762237
```


4.4.2 Task 4b. Show results with a confusion matrix

Use the `confusion_matrix` function to obtain a confusion matrix. Use `y_test` and `y_preds` as arguments.

```
[36]: ### YOUR CODE HERE ###  
cm = confusion_matrix(y_test, y_preds)
```

Next, use the `ConfusionMatrixDisplay()` function to display the confusion matrix from the above cell, passing the confusion matrix you just created as its argument.

```
[37]: ### YOUR CODE HERE ###  
disp = ConfusionMatrixDisplay(confusion_matrix=cm,  
                              display_labels=['retained', 'churned'],  
                              )  
disp.plot();
```



You can use the confusion matrix to compute precision and recall manually. You can also use scikit-learn's `classification_report()` function to generate a table from `y_test` and `y_preds`.

```
[38]: # Calculate precision manually  
precision = cm[1,1] / (cm[0, 1] + cm[1, 1])
```

```
precision
```

```
[38]: 0.5178571428571429
```

```
[39]: # Calculate recall manually
      ### YOUR CODE HERE ###
      # Calculate recall manually
      recall = cm[1,1] / (cm[1, 0] + cm[1, 1])
      recall
```

```
[39]: 0.0914826498422713
```

```
[40]: # Create a classification report
      ### YOUR CODE HERE ###
      # Create a classification report
      target_labels = ['retained', 'churned']
      print(classification_report(y_test, y_preds, target_names=target_labels))
```

	precision	recall	f1-score	support
retained	0.83	0.98	0.90	2941
churned	0.52	0.09	0.16	634
accuracy			0.82	3575
macro avg	0.68	0.54	0.53	3575
weighted avg	0.78	0.82	0.77	3575

Note: The model has decent precision but very low recall, which means that it makes a lot of false negative predictions and fails to capture users who will churn.

4.4.3 BONUS

Generate a bar graph of the model's coefficients for a visual representation of the importance of the model's features.


```
[41]: # Create a list of (column_name, coefficient) tuples
      feature_importance = list(zip(X_train.columns, model.coef_[0]))

      # Sort the list by coefficient value
      feature_importance = sorted(feature_importance, key=lambda x: x[1],
      ↪reverse=True)
      feature_importance
```

```
[41]: [('drives', 0.001913369447769776),
      ('total_navigations_fav1', 0.001231754741616306),
      ('total_navigations_fav2', 0.0009314786513814626),
```

```
(
    'total_sessions', 0.00032707088819142904),
    ('duration_minutes_drives', 0.00010909343558951453),
    ('km_per_driving_day', 1.8223094015325207e-05),
    ('driven_km_drives', -1.4860453424647997e-05),
    ('n_days_after_onboarding', -0.00040647763730561445),
    ('device2', -0.0010412175209008018),
    ('professional_driver', -0.0015285041567402024),
    ('activity_days', -0.10603196504385491)]
```

```
[42]: # Plot the feature importances
      ### YOUR CODE HERE ###
      # Plot the feature importances
      import seaborn as sns
      sns.barplot(x=[x[1] for x in feature_importance],
                  y=[x[0] for x in feature_importance],
                  orient='h')
      plt.title('Feature importance');
```



4.4.4 Task 4c. Conclusion

Now that you've built your regression model, the next step is to share your findings with the Waze leadership team. Consider the following questions as you prepare to write your executive summary. Think about key points you may want to share with the team, and what information is most relevant to the user churn project.

Questions:

1. What variable most influenced the model's prediction? How? Was this surprising?
2. Were there any variables that you expected to be stronger predictors than they were?
3. Why might a variable you thought to be important not be important in the model?
4. Would you recommend that Waze use this model? Why or why not?
5. What could you do to improve this model?
6. What additional features would you like to have to help improve the model?

What variable most influenced the model's prediction? How? Was this surprising?

activity_days was by far the most important feature in the model. It had a negative correlation

Were there any variables that you expected to be stronger predictors than they were?

Yes. In previous EDA, user churn rate increased as the values in km_per_driving_day increased.

Why might a variable you thought to be important not be important in the model?

In a multiple logistic regression model, features can interact with each other and these interactions

Would you recommend that Waze use this model? Why or why not?

It depends. What would the model be used for? If it's used to drive consequential business decisions

What could you do to improve this model?

New features could be engineered to try to generate better predictive signal, as they often do

What additional features would you like to have to help improve the model?

It would be helpful to have drive-level information for each user (such as drive times, geographic

Congratulations! You've completed this lab. However, you may not notice a green check mark next to this item on Coursera's platform. Please continue your progress regardless of the check mark. Just click on the "save" icon at the top of this notebook to ensure your work has been logged.